

Incremental_Capstone_3_EP

April 11, 2024

```
[3]: import numpy as np
import pandas as pd
import json
```

```
aura_df = pd.read_csv("NSMES1988.csv")
aura_df
```

```
[3]:      Unnamed: 0  visits  nvisits  ovisits  novisits  emergency  hospital  \
0              1        5         0         0         0          0          1
1              2        1         0         2         0          2          0
2              3       13         0         0         0          3          3
3              4       16         0         5         0          1          1
4              5        3         0         0         0          0          0
...          ...    ...      ...      ...      ...      ...      ...
4401          4402       11         0         0         0          0          0
4402          4403       12         0         0         0          0          0
4403          4404       10         0        20         0          1          1
4404          4405       16         1         0         0          0          0
4405          4406        0         0         0         0          0          0
```

```
      health  chronic    adl region  age  afam  gender  married  school  \
0    average        2  normal  other  6.9  yes   male     yes      6
1    average        2  normal  other  7.4  no  female     yes     10
2     poor        4  limited  other  6.6  yes  female     no     10
3     poor        2  limited  other  7.6  no   male     yes      3
4    average        2  limited  other  7.9  no  female     yes      6
...      ...    ...      ...      ...      ...      ...      ...
4401  average        0  normal  other  8.4  no  female     yes      8
4402  average        2  normal  other  7.8  no  female     no     11
4403  average        5  normal  other  7.3  no   male     yes     12
4404  average        0  normal  other  6.6  no  female     yes     12
4405  excellent        0  normal  other  7.1  no   male     yes      0
```

```
      income  employed  insurance  medicaid
0    2.881000        yes        yes        no
1    2.747800         no        yes        no
2    0.653200         no         no        yes
3    0.658800         no        yes        no
```

4	0.658800	no	yes	no
...
4401	2.249700	no	yes	no
4402	5.813200	no	yes	no
4403	3.877916	no	yes	no
4404	3.877916	no	yes	no
4405	6.596800	yes	no	no

[4406 rows x 20 columns]

Aura DataFrame Info

```
[4]: aura_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4406 entries, 0 to 4405
Data columns (total 20 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      4406 non-null  int64
1   visits          4406 non-null  int64
2   nvisits         4406 non-null  int64
3   ovisits         4406 non-null  int64
4   novisits        4406 non-null  int64
5   emergency       4406 non-null  int64
6   hospital        4406 non-null  int64
7   health          4406 non-null  object
8   chronic         4406 non-null  int64
9   adl             4406 non-null  object
10  region          4406 non-null  object
11  age             4406 non-null  float64
12  afam            4406 non-null  object
13  gender          4406 non-null  object
14  married         4406 non-null  object
15  school          4406 non-null  int64
16  income          4406 non-null  float64
17  employed        4406 non-null  object
18  insurance       4406 non-null  object
19  medicaid       4406 non-null  object
dtypes: float64(2), int64(9), object(9)
memory usage: 688.6+ KB
```

Aura DataFrame data types:

```
[5]: aura_df.dtypes
```

```
[5]: Unnamed: 0      int64
visits            int64
nvisits           int64
```

```

ovisits      int64
novisits     int64
emergency    int64
hospital     int64
health       object
chronic      int64
adl          object
region       object
age          float64
afam         object
gender       object
married      object
school       int64
income       float64
employed     object
insurance    object
medicaid    object
dtype: object

```

```

[6]: # unnamed 0 column (corrupt data - should be dropped from dataframe)
      # health, region, afam, gender, married data type is object (corrupt data -
      ↪should be category)
      # employed, insurance, medicaid data type is object (corrupt data - should be
      ↪boolean)

```

Aura DataFrame Description

```
[7]: aura_df.age.describe()
```

```

[7]: count      4406.000000
      mean        7.402406
      std         0.633405
      min         6.600000
      25%         6.900000
      50%         7.300000
      75%         7.800000
      max         10.900000
      Name: age, dtype: float64

```

```

[8]: # age range from 6.6 to 10.9 (corrupt data - should be 66 - 109)
      # age data type is float64 (corrupt data - should be uint8)

```

```
[9]: aura_df.income.describe()
```

```

[9]: count      4406.000000
      mean        2.527132
      std         2.924648
      min        -1.012500

```

```

25%      0.912150
50%      1.698150
75%      3.172850
max      54.835100
Name: income, dtype: float64

```

```
[10]: # income range from -1.01 to 54.84 (corrupt data - should be positive value)
      # income data type is float64 (corrupt data - should be float16)
```

```
[11]: aura_df[aura_df.income < 0]
```

```
[11]:
```

	Unnamed: 0	visits	nvisits	ovisits	novisits	emergency	hospital	\
909	910	10	0	0	0	0	0	
910	911	9	2	0	0	0	0	
2592	2593	6	0	0	0	0	0	

	health	chronic	adl	region	age	afam	gender	married	school	\
909	poor	1	normal	other	7.8	no	male	yes	12	
910	average	1	normal	other	7.5	no	female	yes	14	
2592	average	4	normal	west	6.9	no	male	yes	6	

	income	employed	insurance	medicaid
909	-1.0125	no	no	no
910	-1.0125	no	no	no
2592	-0.8180	no	yes	no

```
[12]: aura_df.isnull().sum()
```

```
[12]: Unnamed: 0    0
      visits      0
      nvisits     0
      ovisits     0
      novisits    0
      emergency   0
      hospital    0
      health      0
      chronic     0
      adl         0
      region      0
      age         0
      afam        0
      gender      0
      married     0
      school      0
      income      0
      employed    0
      insurance   0
```

```

medicaid      0
dtype: int64

```

```
[13]: # there is no missing data
```

```
[14]: aura_df.drop('Unnamed: 0', axis=1, inplace=True)
aura_df
```

```
[14]:
```

	visits	nvisits	ovisits	novisits	emergency	hospital	health \
0	5	0	0	0	0	1	average
1	1	0	2	0	2	0	average
2	13	0	0	0	3	3	poor
3	16	0	5	0	1	1	poor
4	3	0	0	0	0	0	average
...
4401	11	0	0	0	0	0	average
4402	12	0	0	0	0	0	average
4403	10	0	20	0	1	1	average
4404	16	1	0	0	0	0	average
4405	0	0	0	0	0	0	excellent

	chronic	adl	region	age	afam	gender	married	school	income \
0	2	normal	other	6.9	yes	male	yes	6	2.881000
1	2	normal	other	7.4	no	female	yes	10	2.747800
2	4	limited	other	6.6	yes	female	no	10	0.653200
3	2	limited	other	7.6	no	male	yes	3	0.658800
4	2	limited	other	7.9	no	female	yes	6	0.658800
...
4401	0	normal	other	8.4	no	female	yes	8	2.249700
4402	2	normal	other	7.8	no	female	no	11	5.813200
4403	5	normal	other	7.3	no	male	yes	12	3.877916
4404	0	normal	other	6.6	no	female	yes	12	3.877916
4405	0	normal	other	7.1	no	male	yes	0	6.596800

	employed	insurance	medicaid
0	yes	yes	no
1	no	yes	no
2	no	no	yes
3	no	yes	no
4	no	yes	no
...
4401	no	yes	no
4402	no	yes	no
4403	no	yes	no
4404	no	yes	no
4405	yes	no	no

[4406 rows x 19 columns]

Fix age:

```
[15]: aura_df['age'] = aura_df['age'] * 10
aura_df['age'] = aura_df['age'].astype('uint8')
aura_df
```

```
[15]:   visits  nvisits  ovisits  novisits  emergency  hospital  health \
0         5        0        0         0          0         1  average
1         1        0        2         0          2         0  average
2        13        0        0         0          3         3    poor
3        16        0        5         0          1         1    poor
4         3        0        0         0          0         0  average
...     ...      ...      ...      ...      ...      ...
4401      11        0        0         0          0         0  average
4402      12        0        0         0          0         0  average
4403      10        0       20         0          1         1  average
4404      16        1        0         0          0         0  average
4405        0        0        0         0          0         0  excellent
```

```
   chronic    adl region  age afam  gender married  school  income \
0         2  normal  other  69  yes   male     yes      6  2.881000
1         2  normal  other  74  no   female   yes     10  2.747800
2         4  limited  other  66  yes  female    no     10  0.653200
3         2  limited  other  76  no   male     yes      3  0.658800
4         2  limited  other  79  no  female     yes      6  0.658800
...     ...      ...      ...      ...      ...      ...
4401      0  normal  other  84  no   female   yes      8  2.249700
4402      2  normal  other  78  no   female    no     11  5.813200
4403      5  normal  other  73  no   male     yes     12  3.877916
4404      0  normal  other  66  no   female   yes     12  3.877916
4405      0  normal  other  71  no   male     yes      0  6.596800
```

```
   employed  insurance  medicaid
0         yes         yes        no
1         no          yes        no
2         no          no         yes
3         no          yes        no
4         no          yes        no
...     ...      ...      ...
4401      no          yes        no
4402      no          yes        no
4403      no          yes        no
4404      no          yes        no
4405      yes         no         no
```

[4406 rows x 19 columns]

```
[16]: # replace yes/no cells with 1/0 so the LLM can actually use the numerical data
```

```
[17]: yes_no = {'yes':1, 'no':0}
yes_no_columns = ['afam', 'married', 'employed', 'insurance', 'medicaid']
aura_df[yes_no_columns] = aura_df[yes_no_columns].replace(yes_no)
aura_df[yes_no_columns] = aura_df[yes_no_columns].astype('uint8')
aura_df
```

```
/var/folders/_8/7k3l29n14j91jpxpd6v9rmtw0000gn/T/ipykernel_7185/3602872080.py:3:
```

```
FutureWarning: Downcasting behavior in `replace` is deprecated and will be
removed in a future version. To retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
```

```
aura_df[yes_no_columns] = aura_df[yes_no_columns].replace(yes_no)
```

```
[17]:
```

	visits	nvisits	ovisits	novisits	emergency	hospital	health \
0	5	0	0	0	0	1	average
1	1	0	2	0	2	0	average
2	13	0	0	0	3	3	poor
3	16	0	5	0	1	1	poor
4	3	0	0	0	0	0	average
...	
4401	11	0	0	0	0	0	average
4402	12	0	0	0	0	0	average
4403	10	0	20	0	1	1	average
4404	16	1	0	0	0	0	average
4405	0	0	0	0	0	0	excellent

	chronic	adl	region	age	afam	gender	married	school	income \
0	2	normal	other	69	1	male	1	6	2.881000
1	2	normal	other	74	0	female	1	10	2.747800
2	4	limited	other	66	1	female	0	10	0.653200
3	2	limited	other	76	0	male	1	3	0.658800
4	2	limited	other	79	0	female	1	6	0.658800
...		
4401	0	normal	other	84	0	female	1	8	2.249700
4402	2	normal	other	78	0	female	0	11	5.813200
4403	5	normal	other	73	0	male	1	12	3.877916
4404	0	normal	other	66	0	female	1	12	3.877916
4405	0	normal	other	71	0	male	1	0	6.596800

	employed	insurance	medicaid
0	1	1	0
1	0	1	0
2	0	0	1
3	0	1	0
4	0	1	0
...

```

4401      0      1      0
4402      0      1      0
4403      0      1      0
4404      0      1      0
4405      1      0      0

```

[4406 rows x 19 columns]

```
[18]: # set columns of category data type
```

```
[19]: category_columns = ['health', 'adl', 'gender', 'region']
aura_df[category_columns] = aura_df[category_columns].astype('category')
aura_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4406 entries, 0 to 4405
Data columns (total 19 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   visits     4406 non-null   int64
 1   nvisits    4406 non-null   int64
 2   ovisits    4406 non-null   int64
 3   novisits   4406 non-null   int64
 4   emergency  4406 non-null   int64
 5   hospital   4406 non-null   int64
 6   health     4406 non-null   category
 7   chronic    4406 non-null   int64
 8   adl        4406 non-null   category
 9   region     4406 non-null   category
10   age        4406 non-null   uint8
11   afam       4406 non-null   uint8
12   gender     4406 non-null   category
13   married    4406 non-null   uint8
14   school     4406 non-null   int64
15   income     4406 non-null   float64
16   employed   4406 non-null   uint8
17   insurance  4406 non-null   uint8
18   medicaid  4406 non-null   uint8
dtypes: category(4), float64(1), int64(8), uint8(6)
memory usage: 353.5 KB

```

Optimizing data by reducing memory size:

```
[20]: int64_columns = aura_df.select_dtypes('int64').columns
aura_df[int64_columns].describe()
```

```

[20]:      visits      nvisits      ovisits      novisits      emergency \
count  4406.000000  4406.000000  4406.000000  4406.000000  4406.000000

```


mean	5.774399	1.618021	0.750794	0.536087	0.263504
std	6.759225	5.317056	3.652759	3.879506	0.703659
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000	0.000000	0.000000
50%	4.000000	0.000000	0.000000	0.000000	0.000000
75%	8.000000	1.000000	0.000000	0.000000	0.000000
max	89.000000	104.000000	141.000000	155.000000	12.000000

	hospital	chronic	school
count	4406.000000	4406.000000	4406.000000
mean	0.295960	1.541988	10.290286
std	0.746398	1.349632	3.738736
min	0.000000	0.000000	0.000000
25%	0.000000	1.000000	8.000000
50%	0.000000	1.000000	11.000000
75%	0.000000	2.000000	12.000000
max	8.000000	8.000000	18.000000

```
[21]: aura_df[int64_columns] = aura_df[int64_columns].astype('int16')
aura_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4406 entries, 0 to 4405
Data columns (total 19 columns):
#   Column      Non-Null Count  Dtype
---  -
0   visits      4406 non-null   int16
1   nvisits     4406 non-null   int16
2   ovisits     4406 non-null   int16
3   novisits    4406 non-null   int16
4   emergency   4406 non-null   int16
5   hospital    4406 non-null   int16
6   health      4406 non-null   category
7   chronic     4406 non-null   int16
8   adl         4406 non-null   category
9   region      4406 non-null   category
10  age         4406 non-null   uint8
11  afam        4406 non-null   uint8
12  gender      4406 non-null   category
13  married     4406 non-null   uint8
14  school      4406 non-null   int16
15  income      4406 non-null   float64
16  employed    4406 non-null   uint8
17  insurance    4406 non-null   uint8
18  medicaid   4406 non-null   uint8
dtypes: category(4), float64(1), int16(8), uint8(6)
memory usage: 147.0 KB
```

```
[22]: float64_columns = aura_df.select_dtypes('float64').columns
aura_df[float64_columns] = aura_df[float64_columns].astype('float16')
aura_df[float64_columns].info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4406 entries, 0 to 4405
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0   income  4406 non-null    float16
dtypes: float16(1)
memory usage: 8.7 KB
```

```
[23]: # fix negative income
aura_df.loc[aura_df.income < 0, 'income'] = 0
aura_df.income.describe()
```

```
[23]: count      4406.000000
mean         2.525391
std          2.921875
min          0.000000
25%          0.912231
50%          1.697754
75%          3.173340
max          54.843750
Name: income, dtype: float64
```

Exported optimized data to a CSV and PKL file:

```
[24]: aura_df.to_csv("NSMES1988_optimized.csv", index=False)
optimized_aura_df = pd.read_csv("NSMES1988_optimized.csv")
optimized_aura_df.info()

# to preserve the data types, export dataframe to a pkl file

aura_df.to_pickle("NSMES1988_optimized.pkl")
optimized_aura_df = pd.read_pickle("NSMES1988_optimized.pkl")
optimized_aura_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4406 entries, 0 to 4405
Data columns (total 19 columns):
#   Column      Non-Null Count  Dtype
---  -
0   visits      4406 non-null    int64
1   nvisits      4406 non-null    int64
2   ovisits      4406 non-null    int64
3   novisits     4406 non-null    int64
4   emergency    4406 non-null    int64
```

```

5  hospital  4406 non-null  int64
6  health    4406 non-null  object
7  chronic   4406 non-null  int64
8  adl       4406 non-null  object
9  region    4406 non-null  object
10 age       4406 non-null  int64
11 afam      4406 non-null  int64
12 gender    4406 non-null  object
13 married   4406 non-null  int64
14 school    4406 non-null  int64
15 income    4406 non-null  float64
16 employed  4406 non-null  int64
17 insurance 4406 non-null  int64
18 medicaid 4406 non-null  int64
dtypes: float64(1), int64(14), object(4)
memory usage: 654.1+ KB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4406 entries, 0 to 4405
Data columns (total 19 columns):
#   Column      Non-Null Count  Dtype
---  -
0   visits      4406 non-null   int16
1   nvisits     4406 non-null   int16
2   ovisits     4406 non-null   int16
3   novisits    4406 non-null   int16
4   emergency   4406 non-null   int16
5   hospital    4406 non-null   int16
6   health      4406 non-null   category
7   chronic     4406 non-null   int16
8   adl         4406 non-null   category
9   region      4406 non-null   category
10  age         4406 non-null   uint8
11  afam        4406 non-null   uint8
12  gender      4406 non-null   category
13  married     4406 non-null   uint8
14  school      4406 non-null   int16
15  income      4406 non-null   float16
16  employed    4406 non-null   uint8
17  insurance   4406 non-null   uint8
18  medicaid   4406 non-null   uint8
dtypes: category(4), float16(1), int16(8), uint8(6)
memory usage: 120.7 KB

```

```

[25]: # Show significant correlations
#aura_df.corr()[abs(aura_df.corr()) > 0.25].fillna(0).style.
      ↪background_gradient(cmap='Spectral', axis=None)

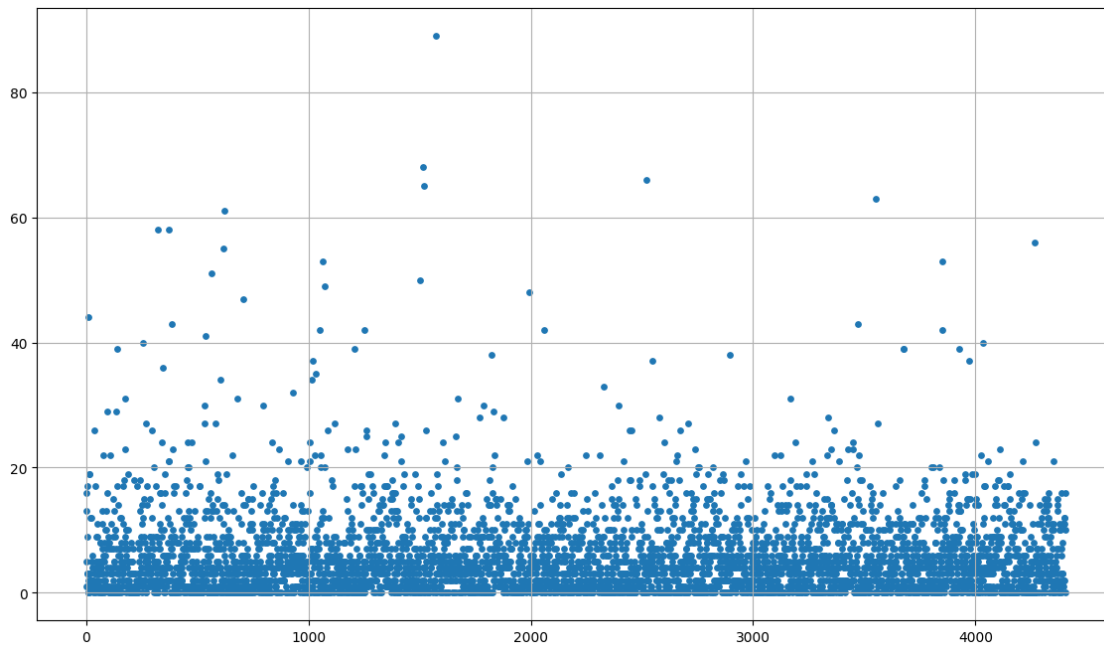
```

```
[26]: # short report detailing visual observations such as number of visits
import matplotlib.pyplot as plt

x = aura_df.index
y = aura_df.visits

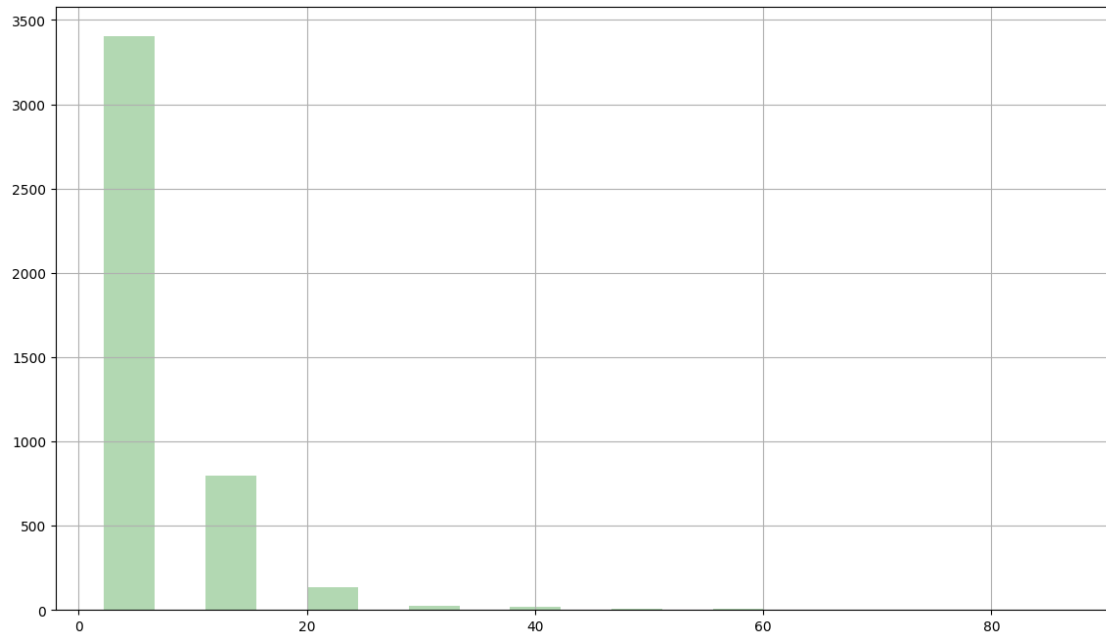
plt.figure(figsize=(14,8))
plt.grid()
plt.scatter(x, y, s=15)
```

```
[26]: <matplotlib.collections.PathCollection at 0x16128ffa0>
```



```
[27]: plt.figure(figsize=(14,8))
plt.grid()
plt.hist(aura_df.visits, bins=10, color='g', alpha=0.3, rwidth=0.5)
```

```
[27]: (array([3.406e+03, 7.970e+02, 1.380e+02, 2.700e+01, 2.100e+01, 7.000e+00,
        5.000e+00, 4.000e+00, 0.000e+00, 1.000e+00]),
array([ 0. ,  8.9, 17.8, 26.7, 35.6, 44.5, 53.4, 62.3, 71.2, 80.1, 89. ]),
<BarContainer object of 10 artists>)
```



```
[28]: optimized_aura_df
```

```
[28]:
```

	visits	nvisits	ovisits	novisits	emergency	hospital	health \
0	5	0	0	0	0	1	average
1	1	0	2	0	2	0	average
2	13	0	0	0	3	3	poor
3	16	0	5	0	1	1	poor
4	3	0	0	0	0	0	average
...	
4401	11	0	0	0	0	0	average
4402	12	0	0	0	0	0	average
4403	10	0	20	0	1	1	average
4404	16	1	0	0	0	0	average
4405	0	0	0	0	0	0	excellent

	chronic	adl	region	age	afam	gender	married	school	income \
0	2	normal	other	69	1	male	1	6	2.880859
1	2	normal	other	74	0	female	1	10	2.748047
2	4	limited	other	66	1	female	0	10	0.653320
3	2	limited	other	76	0	male	1	3	0.658691
4	2	limited	other	79	0	female	1	6	0.658691
...		
4401	0	normal	other	84	0	female	1	8	2.250000
4402	2	normal	other	78	0	female	0	11	5.812500
4403	5	normal	other	73	0	male	1	12	3.876953
4404	0	normal	other	66	0	female	1	12	3.876953

4405	0	normal	other	71	0	male	1	0	6.597656
------	---	--------	-------	----	---	------	---	---	----------

	employed	insurance	medicaid
0	1	1	0
1	0	1	0
2	0	0	1
3	0	1	0
4	0	1	0
...
4401	0	1	0
4402	0	1	0
4403	0	1	0
4404	0	1	0
4405	1	0	0

[4406 rows x 19 columns]

```
[29]: # multiply income by 10 to correct the corrupted data depicting income as X.xx


optimized_aura_df['income'] = optimized_aura_df['income'] * 10
optimized_aura_df['income']
```

```
[29]: 0      28.812500
      1      27.484375
      2       6.531250
      3       6.585938
      4       6.585938
      ...
      4401    22.500000
      4402    58.125000
      4403    38.781250
      4404    38.781250
      4405    66.000000
      Name: income, Length: 4406, dtype: float16
```

```
[30]: df_quant = optimized_aura_df.income.quantile(0.98)
      filtered_data = optimized_aura_df[optimized_aura_df.income < df_quant]

      plt.figure(figsize=(14,8))
      plt.grid()
      plt.scatter(filtered_data.age, filtered_data.income, s=15)
```

```
[30]: <matplotlib.collections.PathCollection at 0x1614113a0>
```



The highest income band occurs under the age of 80.

```
[31]: # visualize how many women and men there are by age groups

grp_women = optimized_aura_df[optimized_aura_df['gender']=='female']
grp_men = optimized_aura_df[optimized_aura_df['gender']=='male']

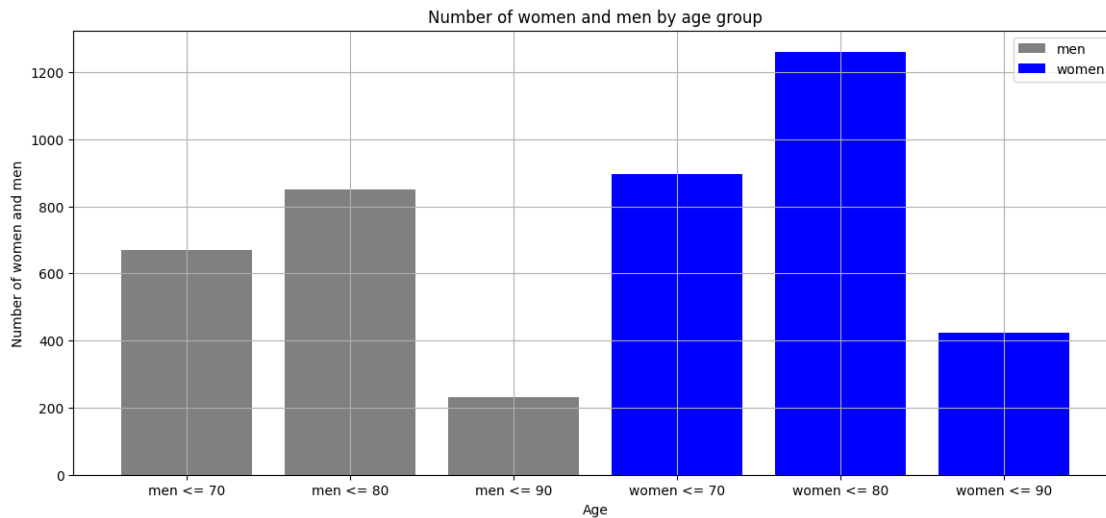
grp_women_70 = grp_women[grp_women['age'] <= 70].shape[0]
grp_men_70 = grp_men[grp_men['age'] <= 70].shape[0]

grp_women_80 = grp_women[(grp_women['age'] > 70) & (grp_women['age'] <= 80)].
↳shape[0]
grp_men_80 = grp_men[(grp_men['age'] > 70) & (grp_men['age'] <= 80)].shape[0]

grp_women_90 = grp_women[(grp_women['age'] > 80) & (grp_women['age'] <= 90)].
↳shape[0]
grp_men_90 = grp_men[(grp_men['age'] > 80) & (grp_men['age'] <= 90)].shape[0]

plt.figure(figsize=(14,6))
plt.grid()
plt.bar(['men <= 70', 'men <= 80', 'men <= 90'],[grp_men_70, grp_men_80,↳
↳grp_men_90], color='gray', label='men')
plt.bar(['women <= 70', 'women <= 80', 'women <= 90'],[grp_women_70,↳
↳grp_women_80, grp_women_90], color='blue', label='women')
plt.xlabel('Age')
plt.ylabel('Number of women and men')
```

```
plt.title('Number of women and men by age group')
plt.legend()
plt.show()
```



Most men and women are between ages 71 and 80. Women between ages 71 and 80 consist of the largest age group by gender. Men between ages 81 and 90 consist of the smallest age group by gender.

```
[32]: optimized_aura_df.age.describe()
```

```
[32]: count    4406.000000
      mean      74.024058
      std       6.334050
      min       66.000000
      25%       69.000000
      50%       73.000000
      75%       78.000000
      max      109.000000
      Name: age, dtype: float64
```

When we describe our dataframe, we confirm most men and women are between ages 71 and 80.

```
[33]: optimized_aura_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4406 entries, 0 to 4405
Data columns (total 19 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   visits      4406 non-null   int16
1   nvisits     4406 non-null   int16
```



```

2  ovisits      4406 non-null  int16
3  novisits     4406 non-null  int16
4  emergency    4406 non-null  int16
5  hospital     4406 non-null  int16
6  health       4406 non-null  category
7  chronic      4406 non-null  int16
8  adl          4406 non-null  category
9  region       4406 non-null  category
10 age          4406 non-null  uint8
11 afam         4406 non-null  uint8
12 gender       4406 non-null  category
13 married      4406 non-null  uint8
14 school       4406 non-null  int16
15 income       4406 non-null  float16
16 employed     4406 non-null  uint8
17 insurance    4406 non-null  uint8
18 medicaid    4406 non-null  uint8
dtypes: category(4), float16(1), int16(8), uint8(6)
memory usage: 120.8 KB

```

All members in the 'novisits' column have a value of 0, and therefore this data is not usable for statistical analysis.

```

[34]: optimized_aura_df.to_csv("NSMES1988_optimized_v2.csv", index=False)
      optimized_aura_df_v2 = pd.read_csv("NSMES1988_optimized_v2.csv")

      optimized_aura_df.to_pickle("NSMES1988_optimized_v2.pkl")
      optimized_aura_df_v2 = pd.read_pickle("NSMES1988_optimized_v2.pkl")
      optimized_aura_df_v2.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4406 entries, 0 to 4405
Data columns (total 19 columns):
#   Column      Non-Null Count  Dtype
---  -
0   visits      4406 non-null  int16
1   nvisits     4406 non-null  int16
2   ovisits     4406 non-null  int16
3   novisits    4406 non-null  int16
4   emergency   4406 non-null  int16
5   hospital    4406 non-null  int16
6   health      4406 non-null  category
7   chronic     4406 non-null  int16
8   adl         4406 non-null  category
9   region      4406 non-null  category
10  age         4406 non-null  uint8
11  afam        4406 non-null  uint8
12  gender      4406 non-null  category
13  married     4406 non-null  uint8

```

```

14 school      4406 non-null   int16
15 income      4406 non-null   float16
16 employed    4406 non-null   uint8
17 insurance    4406 non-null   uint8
18 medicaid    4406 non-null   uint8
dtypes: category(4), float16(1), int16(8), uint8(6)
memory usage: 120.7 KB

```

0.1 Conduct data pivoting, using cross tabulation for example, with pairs of categorical features.

0.2 In this case, let's examine the count of observations for combinations of health per region.


```
[87]: pd.crosstab(optimized_aura_df_v2['health'], optimized_aura_df_v2['region'],
↳ margins=True)
```

```
[87]: region      midwest  northeast  other  west  All
health
average          957          694   1237   621  3509
excellent         90           57    105    91   343
poor             110           86    272    86   554
All              1157          837   1614   798  4406
```

```
[90]: health_region_crosstab = pd.crosstab(optimized_aura_df_v2['health'],
↳ optimized_aura_df_v2['region'], normalize='index').round(2)*100
health_region_crosstab
```

```
[90]: region      midwest  northeast  other  west
health
average          27.0          20.0   35.0  18.0
excellent         26.0          17.0   31.0  27.0
poor              20.0          16.0   49.0  16.0
```

```
[92]: health_region_crosstab.plot(kind='bar', stacked=True, figsize=(8,6));
```



Is a patient's health dependent on the region they live in? Since 'health' and 'region' are categorical features, we can conduct the Chi-Squared Test to see if there is independence between the variables.

H_0 : 'health' is independent of 'region'

H_a : 'health' is highly related to 'region'

```
[48]: from scipy.stats import chi2_contingency

chi2, p_value, dof, expected = chi2_contingency(health_region_crosstab)
# Display the p-value in a legible manner
print(f"The p-value is: {p_value:.16f}")
```

The p-value is: 0.0000000000407694


Since the 'p value' is < 0.05 , we reject the null hypothesis. We conclude that a patient's health is highly correlated with the region they live in.

0.3 Create a distribution table that categorizes individuals by their health status, differentiated by gender.

```
[59]: contingency_table = pd.crosstab(optimized_aura_df_v2['health'],  
    ↪ optimized_aura_df_v2['gender'])  
contingency_table  
  
health_gender_distribution = contingency_table.div(contingency_table.  
    ↪ sum(axis=1), axis=0)  
health_gender_distribution.round(2)
```

```
[59]: gender    female  male  
health  
average      0.60  0.40  
excellent    0.56  0.44  
poor         0.62  0.38
```

```
[100]: health_gender_distribution.plot.bar(stacked='True', figsize=(8,5));
```



0.4 Determine if gender and health are independent categorical variables.

H_0 : 'health' and 'gender' are independent of each other.

H_a : 'health' and 'gender' are dependent of each other.

```
[62]: chi2, p_value, dof, expected = chi2_contingency(health_gender_distribution)
      p_value
```

```
[62]: 0.9968561005885185
```

The p value is 0.99, which is > 0.05 .

We therefore accept the null hypothesis, H_0 :, and conclude that 'health' and 'gender' are independent of one another.


0.5 Compile a table to examine the income distribution across genders.

```
[117]: optimized_aura_df_v2['income'] = optimized_aura_df_v2['income'].
      ↪astype('float32') #crosstab does not accept float16
      income_gender_contingency_table = pd.crosstab(optimized_aura_df_v2['income'],
      ↪optimized_aura_df_v2['gender'])
      income_gender_contingency_table
```

```
[117]: gender  female  male
      income
0.0         20     10
1.0          6      1
2.0         21      2
3.0         60     17
4.0        193     33
...
230.0         1      2
242.0         1      0
360.0         0      1
417.0         1      1
548.0         1      1
```

[146 rows x 2 columns]

```
[109]: import seaborn as sns
      sns.boxplot(data=optimized_aura_df_v2, x="gender", y="income");
```



```
[111]: optimized_aura_df_v2.groupby('gender', observed=True)['income'].describe().
       ↪round(2).T
```

```
[111]: gender  female    male
count   2628.00  1778.00
mean     22.01   28.90
std      27.21   31.57
min       0.00    0.00
25%       7.00   12.00
50%      14.00   20.00
75%      27.00   35.00
max     548.00  548.00
```

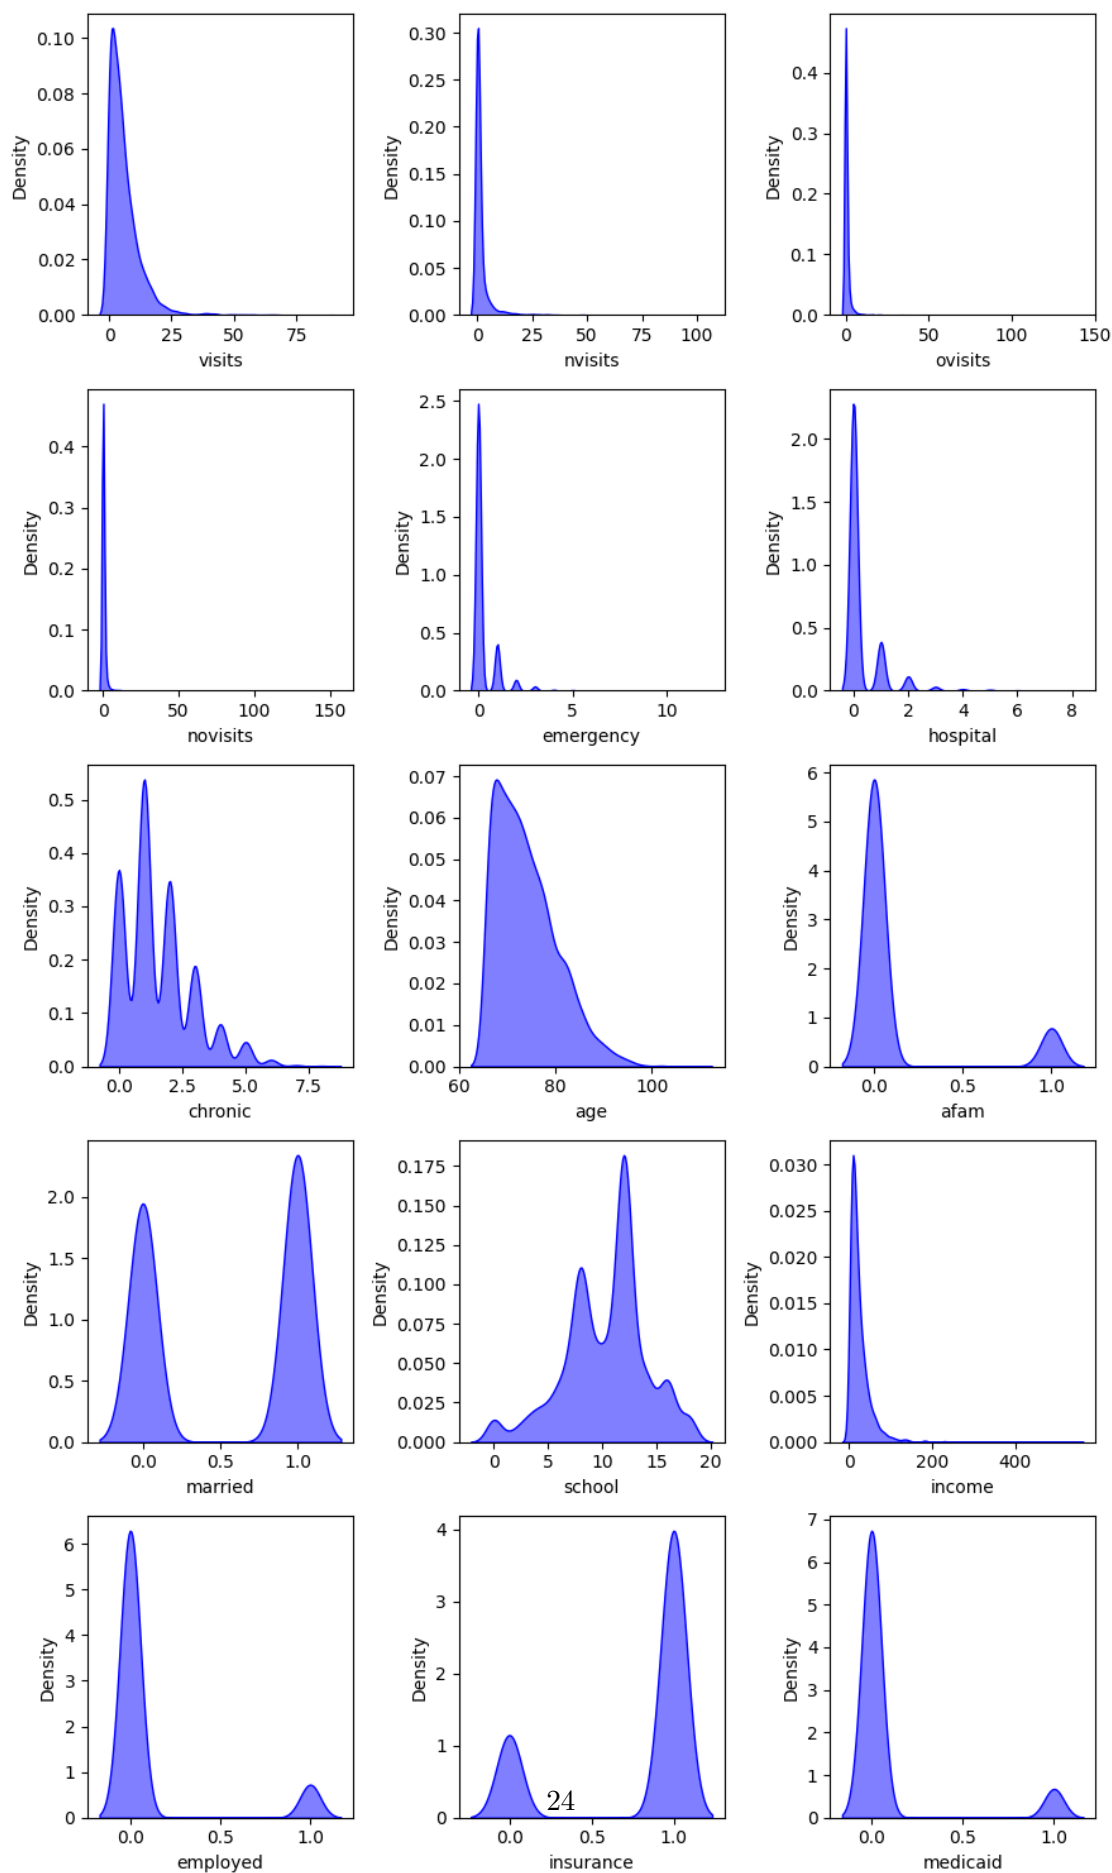
```
[77]: chi2, p_value, dof, expected = chi2_contingency(income_gender_contingency_table)
       print(f"The p-value is: {p_value:.16f}")
```

The p-value is: 0.00000000000003508


Income distribution and gender are highly dependent of one another.

0.6 Univariate Analysis

```
[84]: from vizad.univariate import plot_univariate_numeric,   
      ↪ plot_univariate_categorical   
   
 num_cols = optimized_aura_df_v2.select_dtypes(include=np.number).columns.  
      ↪ tolist()   
 cat_cols = [col for col in optimized_aura_df_v2.columns if col not in num_cols]   
   
 plot_univariate_numeric(optimized_aura_df_v2, num_cols, figsize=(8,8),   
      ↪ kind='density')
```





```
[86]: plot_univariate_categorical(optimized_aura_df_v2, cat_cols, figsize=(8,4))
```



0.7 Bivariate Analysis of Categorical Features

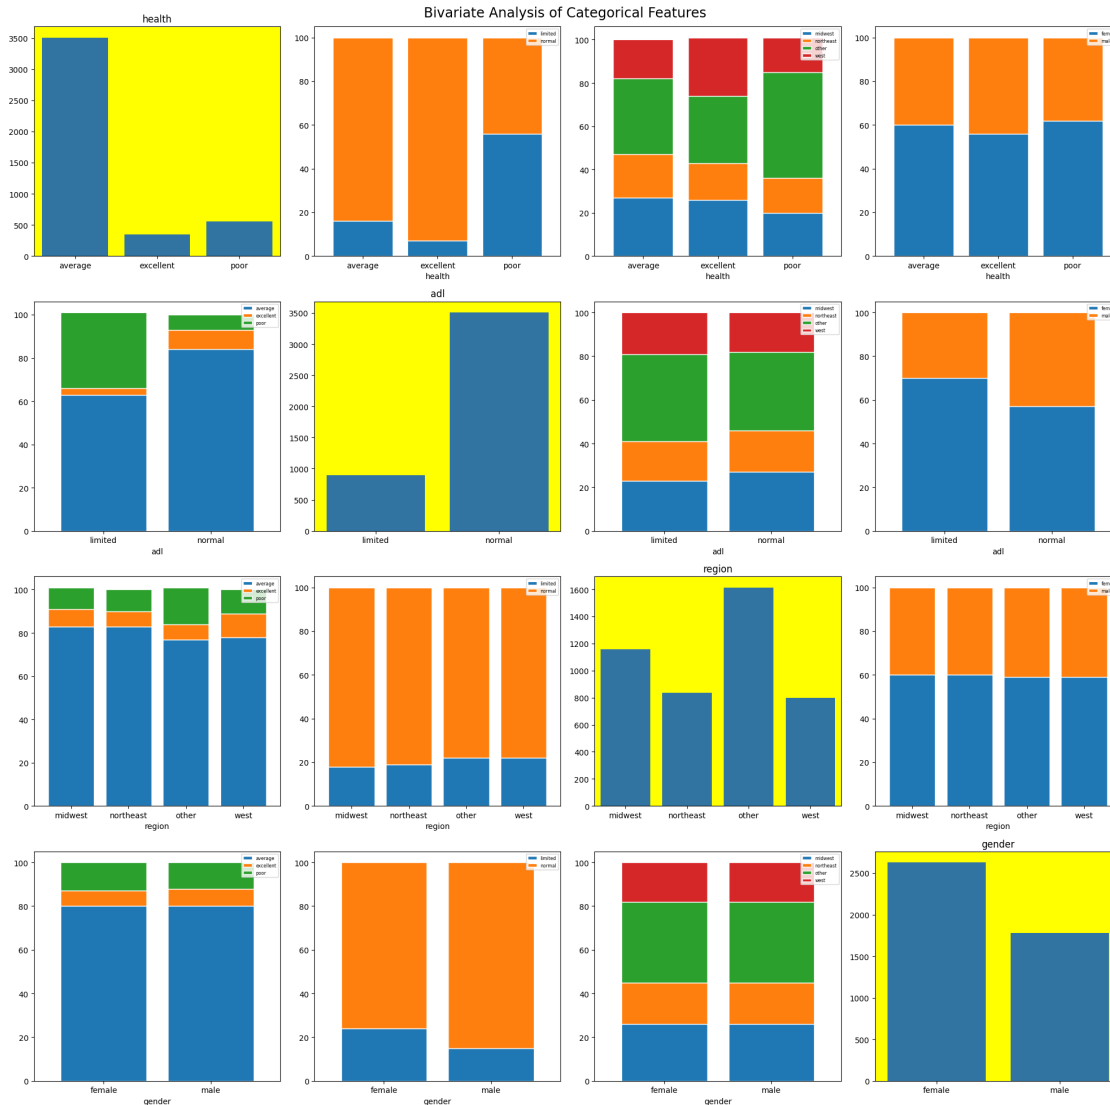
```
[97]: plot_bivariate_numeric(optimized_aura_df_v2, num_cols, kind='heatmap')
```



```
[98]: from vizad.bivariate import plot_bivariate_numeric, plot_bivariate_categorical

num_cols = optimized_aura_df_v2.select_dtypes(include=np.number).columns.
        tolist()
cat_cols = [col for col in optimized_aura_df_v2.columns if col not in num_cols]


plot_bivariate_categorical(optimized_aura_df_v2, cat_cols, figsize=(20,20))
```



0.8 Age vs Income Relationship

Develop a table to analyze the relationship between age and income.

```
[115]: sns.scatterplot(data=optimized_aura_df_v2 , x="age", y="income")
plt.xlabel='age'
plt.ylabel='income'
```



[]: